



Daemons

Daemons

- Understanding Daemons
- Creating a Daemon
- Communicating with a Daemon

Introduction

What is a Daemon?

A **Daemon** is a background process that runs without user input and usually provides some service, either for the system as a whole or for user programs.

Common daemon programs include:

- System Logger (**syslogd**);
- Web Servers (**httpd**);
- Mail Servers (**smtpd**);
- Database Servers (**mysqld**).
- FTP Servers (**ftpd**)

Introduction

Understanding Daemons

- Normally, daemons are started when a system boots and, unless forcibly terminated, run until system shutdown.
- Because a daemon does not have a controlling terminal, any output, either to **stderr** or **stdout**, requires special handling.
- Daemons often run with superuser privilege because they use privileged ports (1 - 1024) or because they have access to some sort of privileged resource.
- Daemons generally are process group leaders and session leaders.
- Finally, a daemon's parent is the **init** process, which has the PID of 1. (daemon is an orphan process inherited by **init**)



Creating a Daemon

Creating a Daemon

Although a daemon may seem mysterious and difficult to program, it is really very simple if you keep a few rules in mind and know the key function calls to make.

Simple steps to follow to create a daemon:

1. **fork** and exit in the parent process;
2. Create a new session in the child using the **setsid** call;
3. Make the root directory, **/**, the child process's working directory;
4. Change the child process's umask to **0**;
5. Close any unneeded file descriptor the child inherited.

Creating a Daemon

1 . Fork and exit in the parent process

A daemon is started from a shell script or the command line.

Daemons are unlike application programs because they are not interactive, i.e., they run in the background and, as a result, do not have the controlling terminal.

The parent forks and exits as the first step toward getting rid of the controlling terminal (they only need a terminal interface long enough to get started).

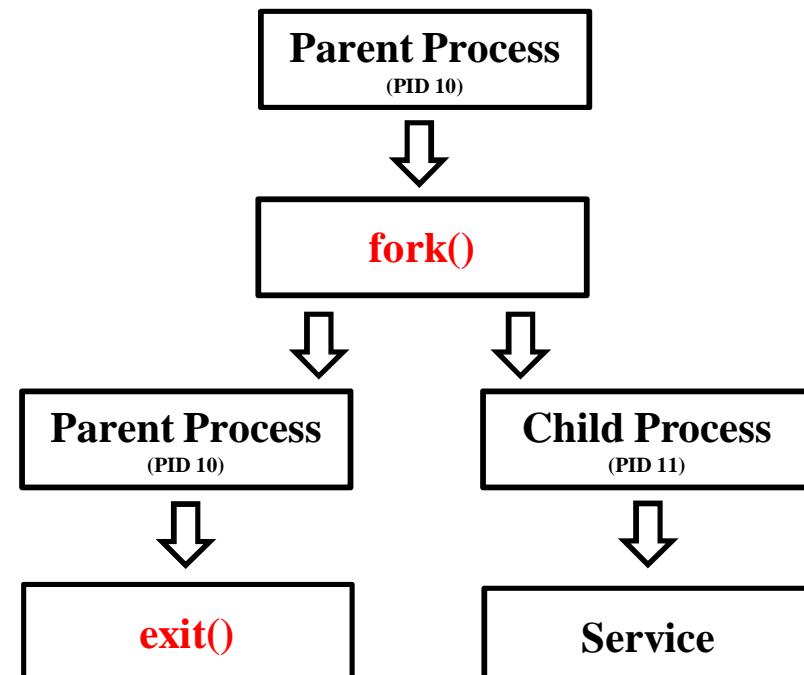
```
pid_t pid;

// create a new process
pid = fork();
if (pid < 0) { // error trying to execute fork
    ERROR("fork failure");
    exit(EXIT_FAILURE);
}

if (pid > 0) // parent process (exit)
exit(EXIT_SUCCESS);
// child process continues the execution
...
```

Creating a Daemon

1 . Fork and exit in the parent process



Creating a Daemon

2 . Create a new session in the child

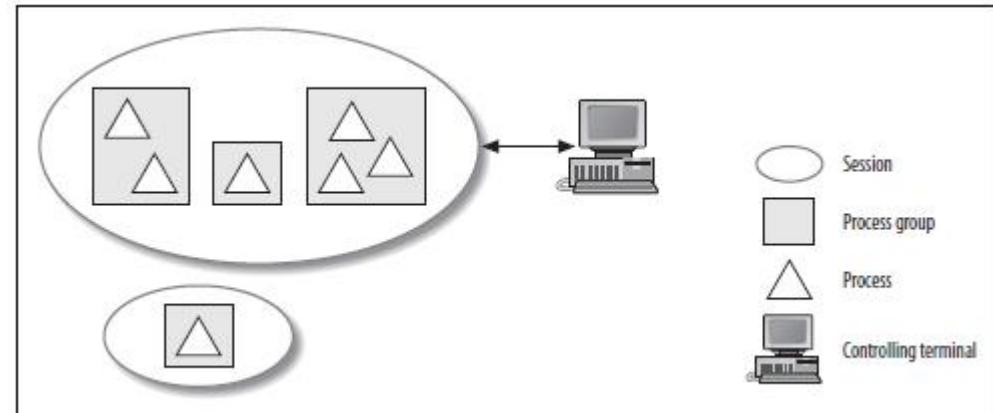
Calling **setsid** accomplishes several things:

- It creates a new session if the calling process is not a process group leader, making the calling process the session leader of the new session;
- It makes the calling process the process group leader of the new process group;
- It sets the process group ID (PGID) and the session ID (SID) to the process ID (PID) of the calling process;
- It dissociates the new session from any controlling **tty**.

```
pid_t sid;
sid = setsid(); // create a new session
if (sid < 0) {
    ERROR("setsid failure");
    exit(EXIT_FAILURE);
}
```

Creating a Daemon

2 . Create a new session in the child



- Each process is a member of a process group, which is a collection of one or more processes generally associated with each other for the purposes of job control (`# cat ship-inventory.txt | grep booty | sort`).
- When a new user first logs into a machine, the login process creates a new session that consists of a single process, the user's login shell. The login shell functions as the session leader.

Creating a Daemon

3 . Make “/” the root directory

This is necessary because any process whose current directory is on a mounted file system will prevent that file system from being unmounted.

Making “/” a daemon’s working directory is a safe way to avoid this possibility.

```
// make '/' the root directory
if (chdir("/") < 0) {
    ERROR("chdir failure");
    exit(EXIT_FAILURE);
}

// continue executing the child process
...
```

Creating a Daemon

4 . Change the child process's umask to 0

This step is necessary to prevent the daemon's inherited umask from interfering with the creation of files and directories.

Consider the following scenario:

- A daemon inherits a umask of 055, which masks out read and execute permissions for group and other.
- Resetting the daemon's umask to 0 prevents such situation.

```
// resetting umask to 0
umask(0);

// continue executing the child process

...
```

Creating a Daemon

5 . Close any unneeded file descriptor

This is simply a common sense step.

There is no reason for a child to keep open descriptors inherited from parent.

The list of potential file descriptors to close includes at least stdin, stdout and stderr.

```
// close unneeded file descriptors
close(STDIN_FILENO);
close(STDOUT_FILENO);
close(STDERR_FILENO);

// continue executing the child process

...
```



Example 1

```
...
int main(int argc, char *argv[])
{
    pid_t pid, sid;

    pid = fork(); // create a new process
    if (pid < 0) { // on error exit
        perror("fork");
        exit(EXIT_FAILURE);
    }

    if (pid > 0) // parent process (exit)
        exit(EXIT_SUCCESS);

    sid = setsid(); // create a new session
    if (sid < 0) { // on error exit
        perror("setsid");
        exit(EXIT_FAILURE);
    }

    if (chdir("/") < 0) { // on error exit
        perror("chdir");
        exit(EXIT_FAILURE);
    }

    umask(0); // make '/' the root directory
    close(STDIN_FILENO); // close standard input file descriptor
    close(STDOUT_FILENO); // close standard output file descriptor
    close(STDERR_FILENO); // close standard error file descriptor
    // service implementation
}
```



Example 1

```
int main(int argc, char *argv[])
{
    ...
    // service implementation
    len = strlen(ctime(&timebuf));
    while (1) {
        char *buf = malloc(sizeof(char) + len + 1);
        if (buf == NULL) {
            perror("malloc");
            exit(EXIT_FAILURE);
        }

        if ((fd = open("/var/log/example1.log",
                      O_CREAT | O_WRONLY | O_APPEND, 0600)) < 0) {
            perror("open");
            exit(EXIT_FAILURE);
        }

        time(&timebuf);
        strncpy(buf, ctime(&timebuf), len + 1);
        write(fd, buf, len + 1);
        close(fd);
        sleep(60);
    }

    exit(EXIT_SUCCESS);
}
```



Example 1

Using example1

Compile example1:

```
# gcc example1.c -o example1
```

Run example1:

```
# sudo ./example1
```

Stop the program:

```
# kill {PID of example1}
```

Open the log file /var/log/example1.log:

```
# sudo tail -f /var/log/example1.log
```



Daemon Problem

Problem:

Once a daemon calls **setsid**, it no longer has the controlling terminal and so it has nowhere to send output that would normally go to **stdout** or **stderr** (such as error messages).

Solution:

Fortunately, the standard utility for this purpose is the **syslog** service, provided by the system logging daemon, **syslogd**.

Using syslog

Handling Errors with syslog

syslogd is a daemon that allow to save log messages from other daemons or applications.

The relevant interface is defined in **<syslog.h>** header file.

The API is simple, **openlog** opens the log, **syslog** writes a message to it, and **closelog** close the log.

The function prototypes are listed here:

```
#include <syslog.h>

void openlog(char *ident, int option, int facility);
void closelog(void);
void syslog(int priority, char *format, ...);
```

Using syslog

```
void openlog(char *ident, int option, int facility);
```

openlog initiates a connection to the system logger.

ident is a string added to each message and typically is set to the name of the program.

The **option** argument is a logical OR of one or more of the default values listed here:

- **LOG_CONS** - write to the console if the system logger is unavailable.
- **LOG_NDELAY** - open the connection immediately. Normally, the connection is not opened until the first message is sent.
- **LOG_PERROR** - print to stderr.
- **LOG_PID** - include the process's PID with each message.

Using syslog

```
void openlog(char *ident, int option, int facility);
```

facility specifies the type of program sending the message and can be one of the following values.

LOG_AUTHPRIV

Security/authorization messages

LOG_CRON

Clock daemons; cron and at

LOG_DAEMON

Other system daemons

LOG_KERN

Kernel messages

LOG_LOCAL[0-7]

Reserved for local use

LOG_LPR

Line printer subsystem

LOG_MAIL

The mail subsystem

LOG_NEWS

Usenet news subsystem

LOG_SYSLOG

Messages generated by syslog

LOG_USER

Default

LOG_UUCP

UUCP

Using syslog

void closelog(void)

closelog close the system logger connection.

The use of **openlog** and **closelog** is optional, because **syslog** will open the log file automatically the first time it is called, and close the log file when your application is terminated.

void syslog(int priority, char *format, ...);

syslog write a message to system logger.

priority specifies the importance of the message.

LOG_EMERG – System is unusable

LOG_ALERT – Take action immediately

LOG_CRIT – Critical condition

LOG_ERR – Error condition

LOG_WARNING – Warning condition

LOG_NOTICE – Normal but significant condition

LOG_INFO – Informational message

LOG_DEBUG – Debugging message



Example 2

```
...
int main(int argc, char *argv[])
{
    pid_t pid, sid;

    pid = fork(); // create a new process
    if (pid < 0) { // on error exit
        syslog(LOG_ERR, "%s\n", "fork");
        exit(EXIT_FAILURE);
    }

    if (pid > 0) // parent process (exit)
        exit(EXIT_SUCCESS);

    sid = setsid(); // create a new session
    if (sid < 0) { // on error exit
        syslog(LOG_ERR, "%s\n", "setsid");
        exit(EXIT_FAILURE);
    }

    if (chdir("/") < 0) { // on error exit
        syslog(LOG_ERR, "%s\n", "chdir");
        exit(EXIT_FAILURE);
    }

    umask(0); // make '/' the root directory
    close(STDIN_FILENO); // close standard input file descriptor
    close(STDOUT_FILENO); // close standard output file descriptor
    close(STDERR_FILENO); // close standard error file descriptor
    // service implementation
}
```



Example 2

Using example2

Compile example2:

```
# gcc example2.c -o example2
```

Run example2:

```
# sudo ./example2
```

Stop the program:

```
# kill {PID of example2}
```

Open the log file /var/log/example1.log:

```
# sudo tail -f /var/log/example2.log
```

Open the syslog:

```
# sudo tail -f /var/log/syslog (messages)
```

Communications

Communicating with a Daemon

To communicate with a daemon, you send it signals that cause it to respond in a given way.

For example, it is typically necessary to force a daemon to reread its configuration file. The most common way to do this is to send a SIGHUP signal to the daemon.

When you execute the command “kill PID” on command line, the signal SIGINT is sent to daemon to terminate the daemon execution.

```
#include <syslog.h>

void openlog(char *ident, int option, int facility);
void closelog(void);
void syslog(int priority, char *format, ...);
```

Signal Example

Communicating with a Daemon

To communicate with a daemon, you send it signals that cause it to respond in a given way.

For example, it is typically necessary to force a daemon to reread its configuration file. The most common way to do this is to send a SIGHUP signal to the daemon.

When you execute the command “kill PID” on command line, the signal SIGINT is sent to daemon to terminate the daemon execution.



Example 3

```
#include <signal.h>

void signal_handler(int sig) {
    switch(sig) {
        case SIGHUP:
            syslog(LOG_INFO,"hangup signal catched");
            break;
        case SIGTERM:
            syslog(LOG_INFO,"terminate signal catched");
            exit(0);
            break;
    }
}

int main(int argc, char *argv[])
{
    pid_t pid, sid;

    pid = fork(); // create a new process
    if (pid < 0) { // on error exit
        syslog(LOG_ERR, "%s\n", "fork");
        exit(EXIT_FAILURE);
    }

    ...

    signal(SIGHUP,signal_handler); /* catch hangup signal */
    signal(SIGTERM,signal_handler); /* catch kill signal */
    // service implementation
}
```



Example 3

Using example3

Compile example3:

```
# gcc example3.c -o example3
```

Run example3:

```
# sudo ./example3
```

Stop the program:

```
# kill {PID of example3}
```

Send SIGHUP signal:

```
# kill -HUP {PID of example3}
```

Open the log file /var/log/example1.log:

```
# sudo tail -f /var/log/example2.log
```

Open the syslog:

```
# sudo tail -f /var/log/syslog (messages)
```